

*Dr. Denton Bobeldyk*

---

# CIS 365 Artificial Intelligence

Perceptron Learning

---



# Week in Review

Blackboard Check-in



# **Delivery Methods**

Lecture

Videos

Lab Time

Small Groups



---

# Decision Making

---

- Multiple factors weigh into determining the outcome of a decision
- Multiple factors weigh into determining the action to take



---

# Decision Making

---

- Would I like to date this person?



---

# Decision Making - Class Exercise

---

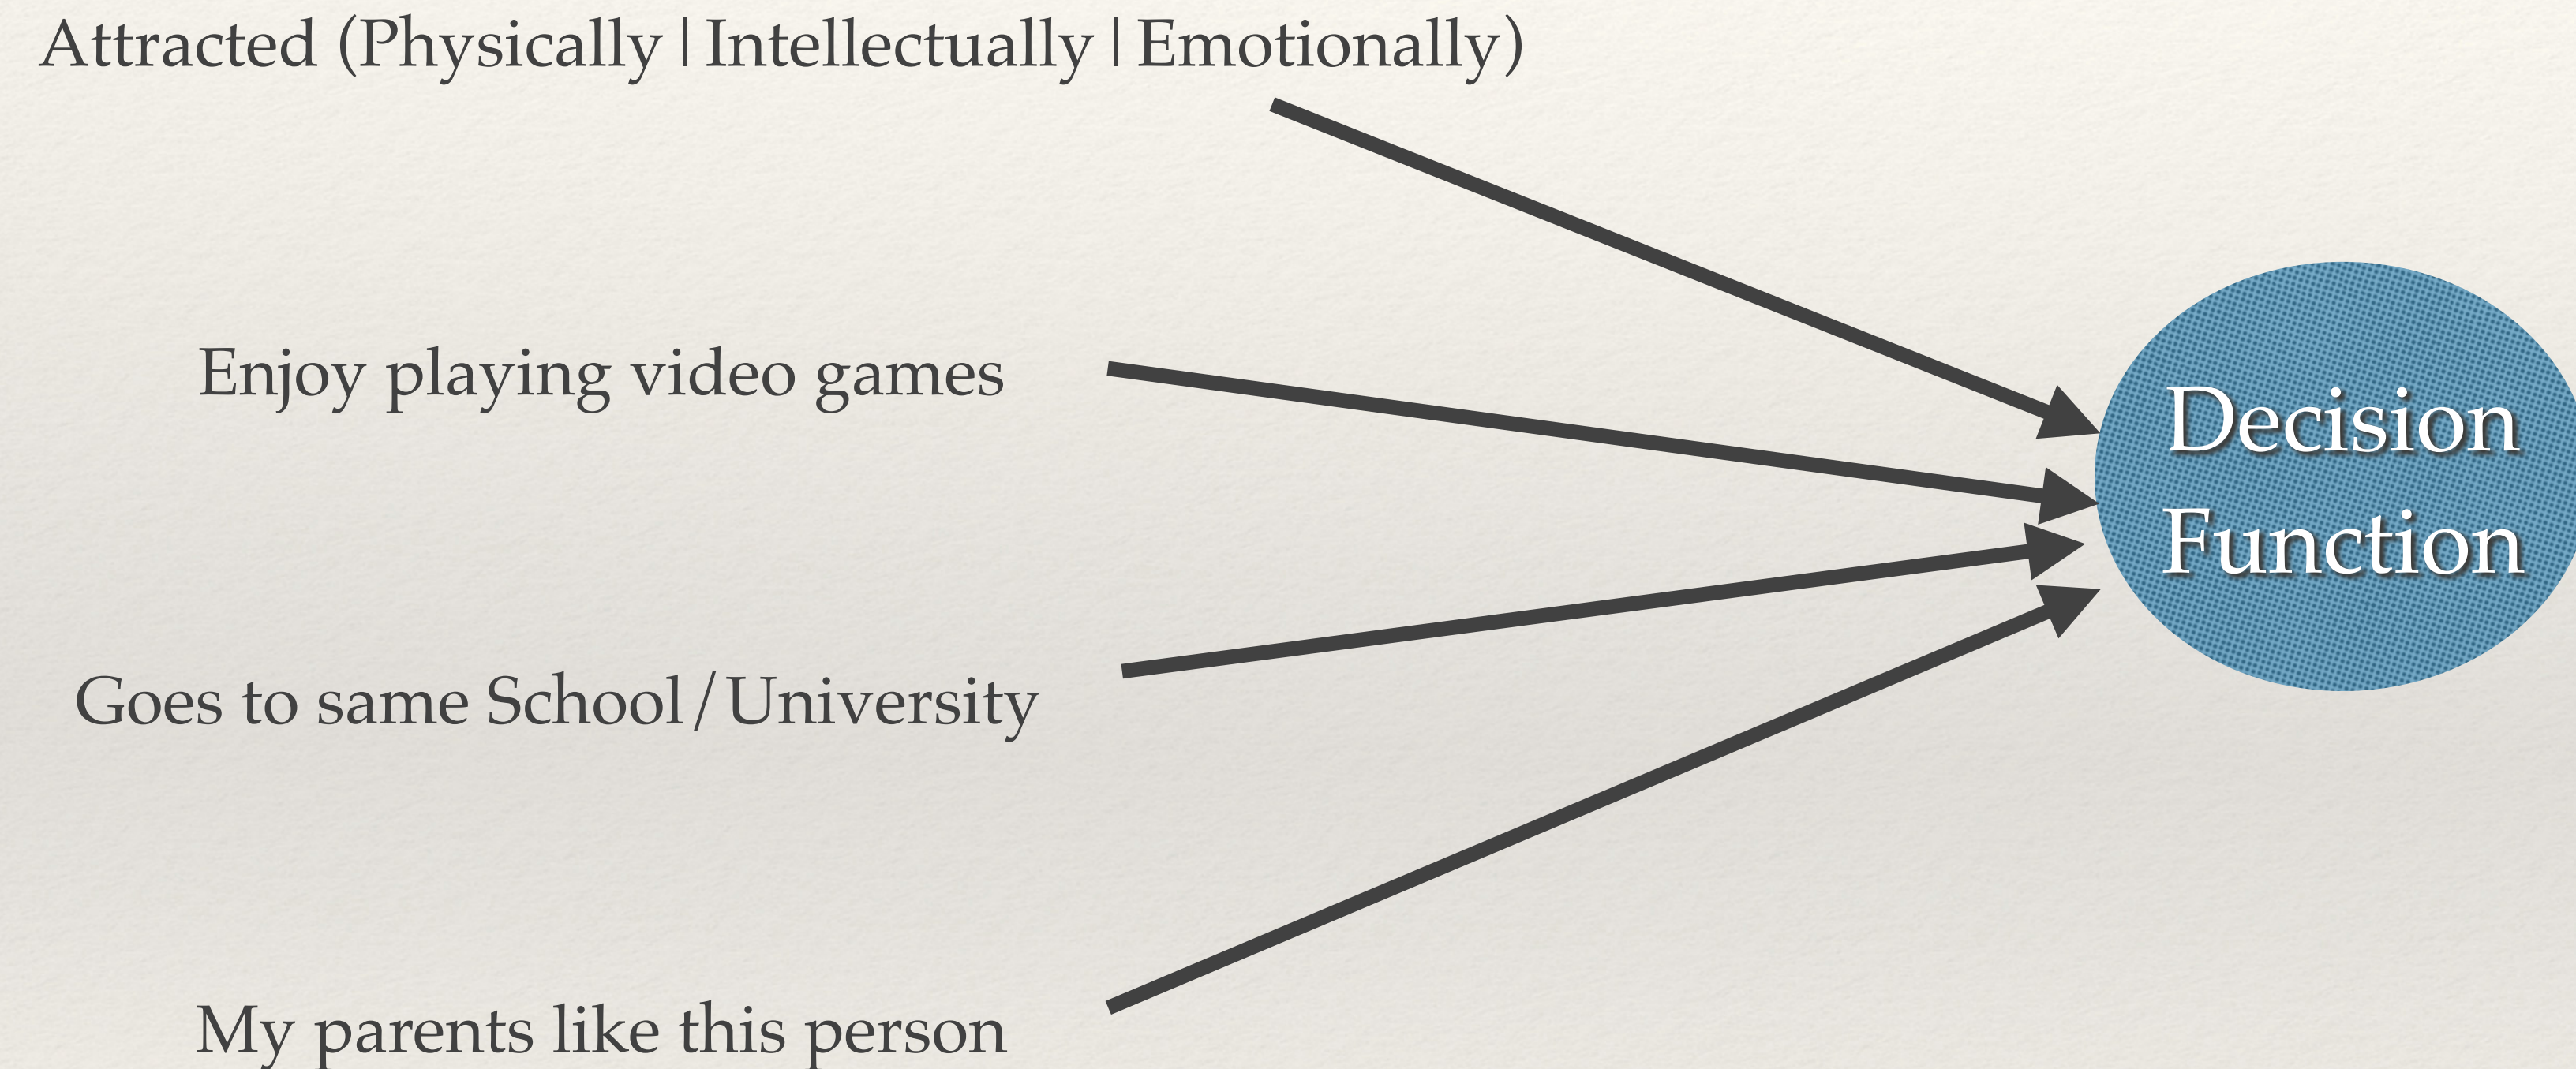
- Would I like to date this person?
  - What considerations go into this decision?



---

# Decision Making

---





---

# Decision Making

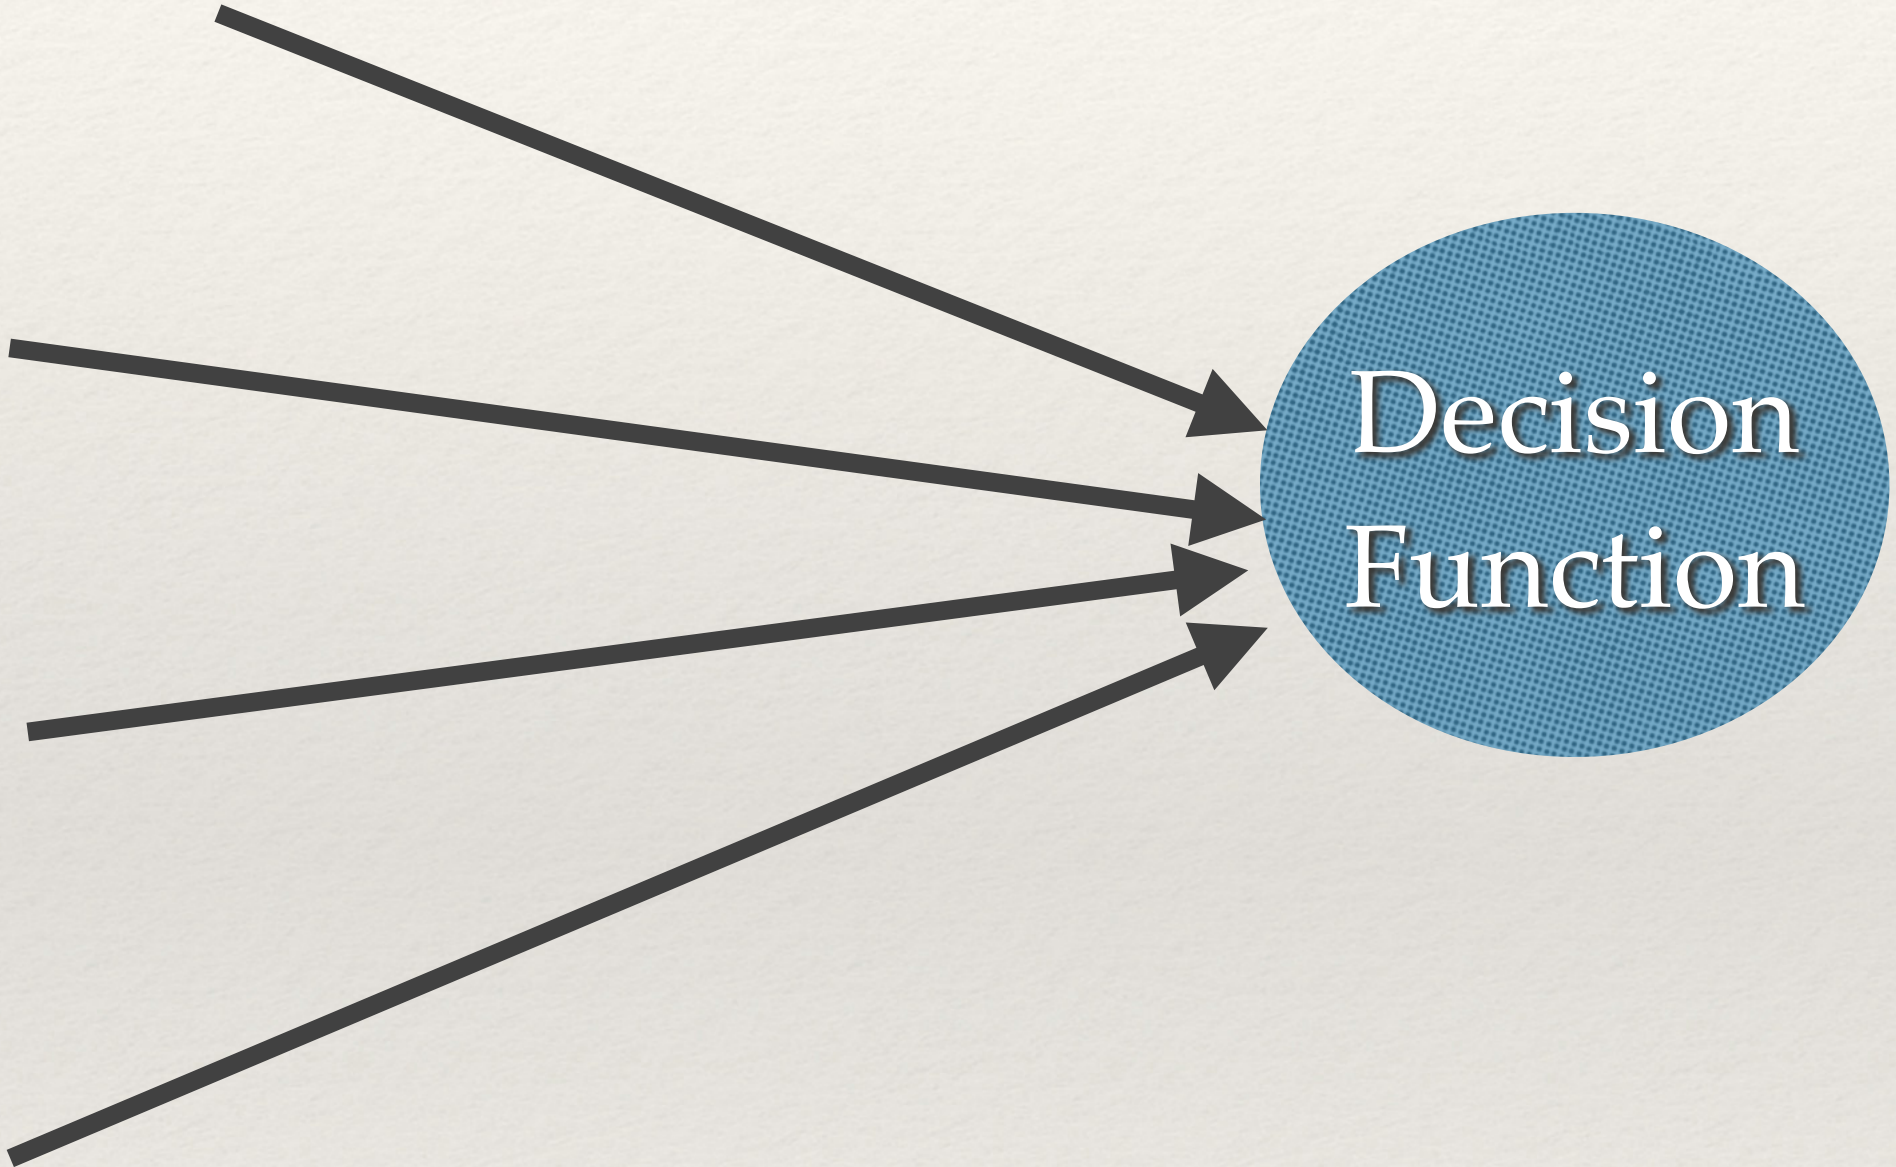
---

Attracted (Physically | Intellectually | Emotionally)

Enjoy playing video games

Goes to same School/University

My parents like this person



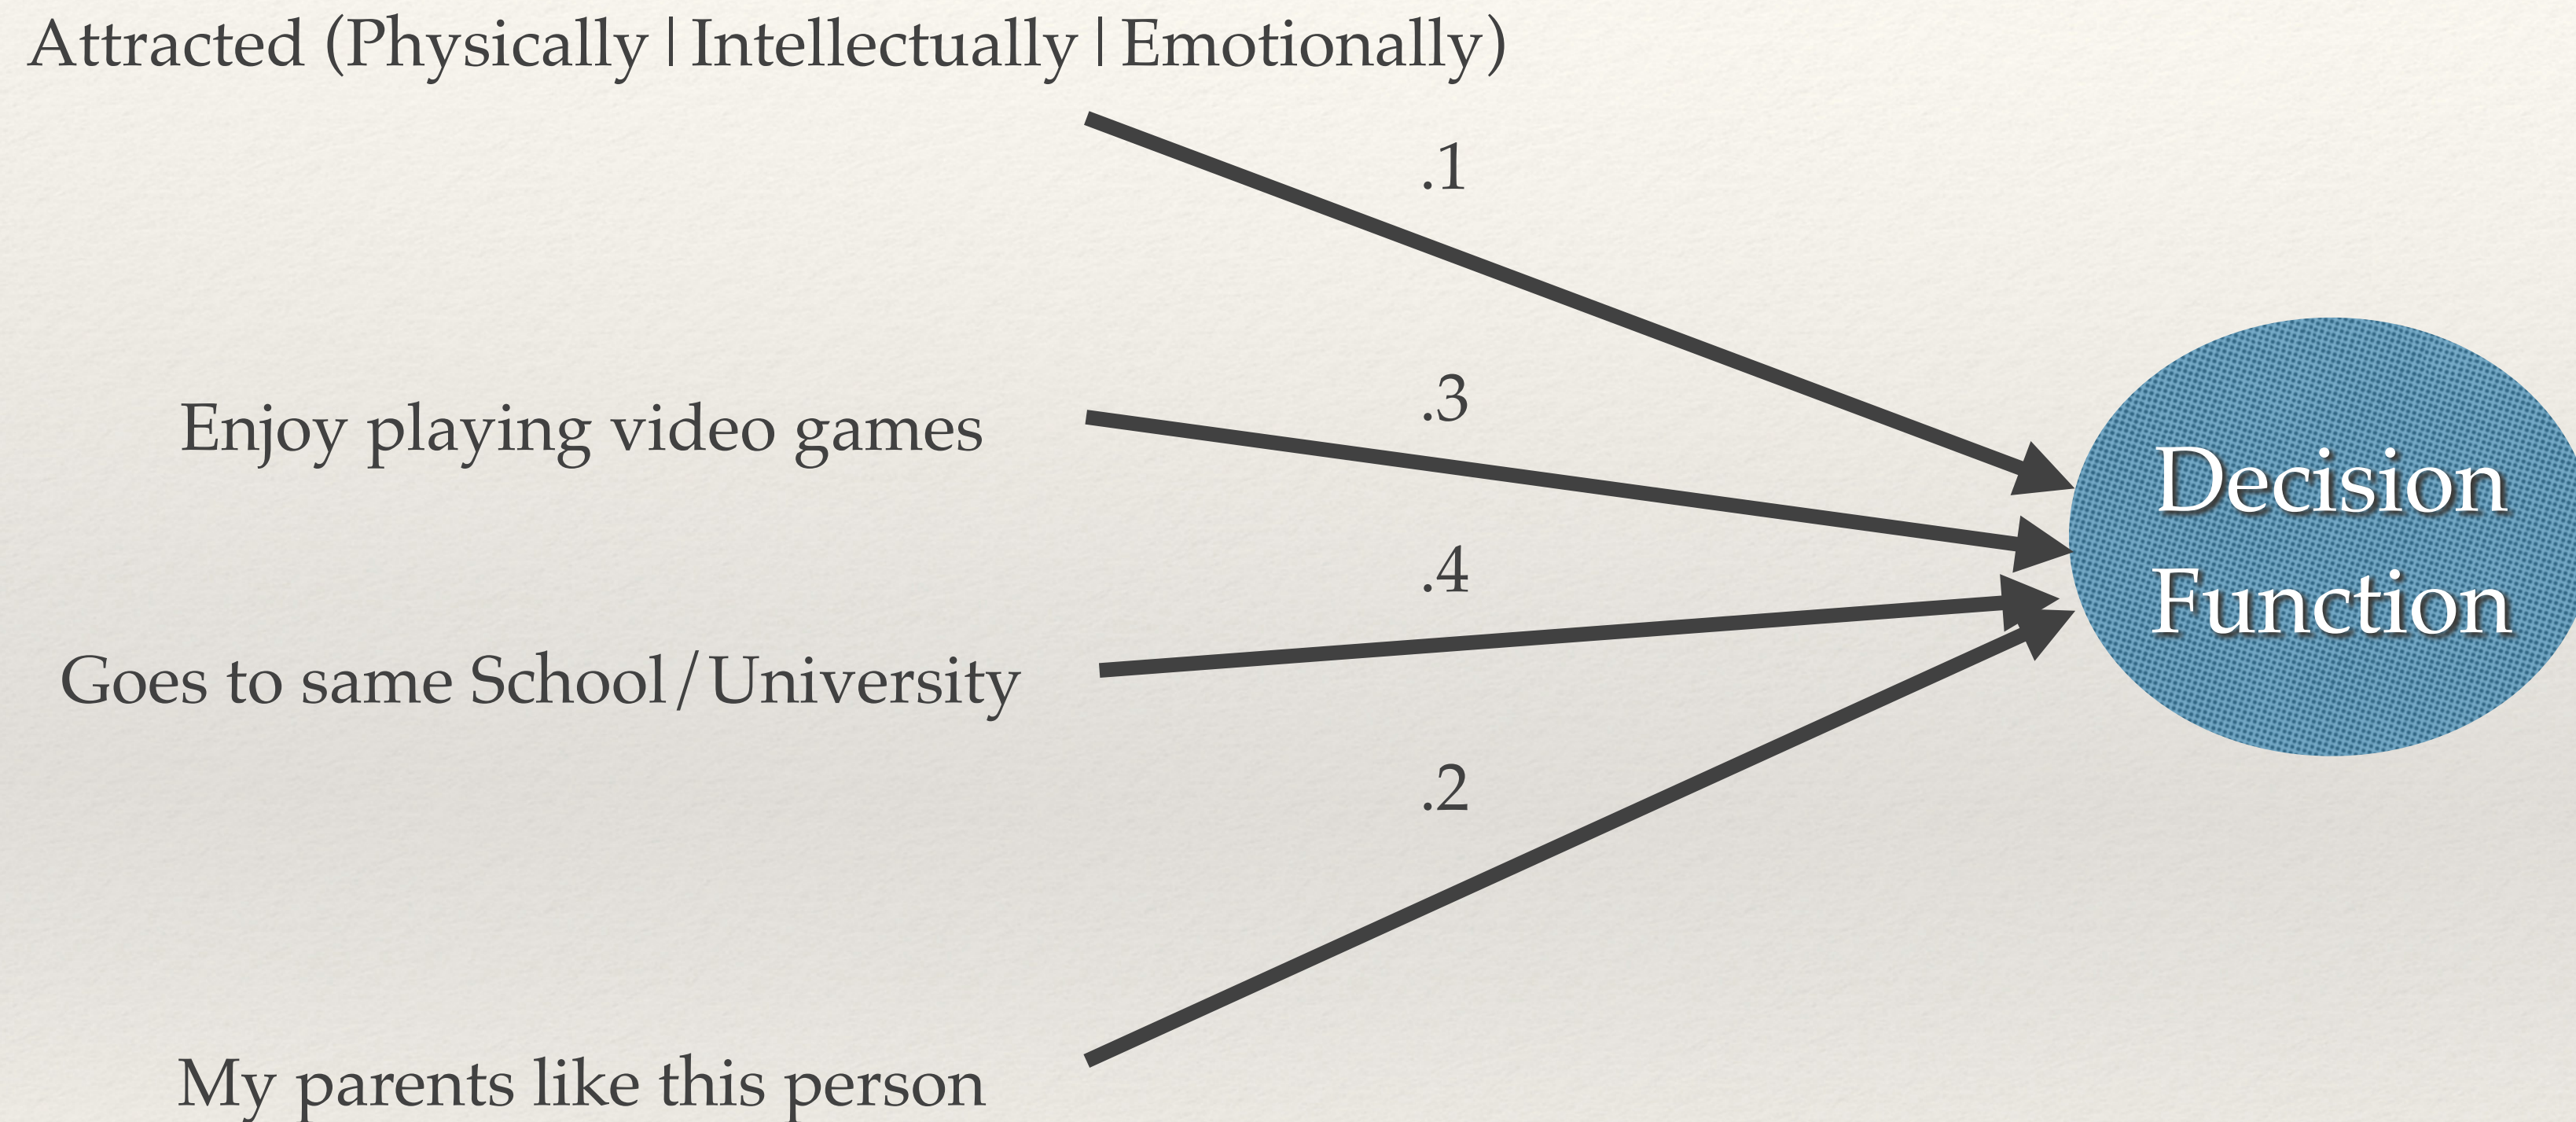
The diagram illustrates a decision-making process. On the left, four text labels are listed vertically: 'Attracted (Physically | Intellectually | Emotionally)', 'Enjoy playing video games', 'Goes to same School/University', and 'My parents like this person'. From each label, a black arrow points diagonally towards a central blue circle on the right. The circle has a textured, dotted pattern and contains the text 'Decision Function' in white. The arrows converge on the left side of the circle, pointing towards its center.

Decision  
Function

Can I weigh how much each of these factor into the decision?




# Decision Making






# Neuron





# Mc Culloch-Pitts Neuron




$$g = \bar{x} \cdot \bar{w}$$

F is an activation function



# Perceptron - General Notation




F is an activation function



# Perceptron with Thresholding




# Perceptron



An abstract neuron referred to as a Threshold Logic Unit



# Perceptron




An abstract neuron referred to as a Threshold Logic Unit



# Threshold Logic Unit

$$\bar{x} = (1, 1, 0)$$

$$\bar{w} = (0.5, 1.0, 1.2)$$




If the output greater than or equal to the threshold send output of 1

$$\bar{x} \cdot \bar{w} = x_1 * w_1 + x_2 * w_2 + x_3 * w_3 =$$

$$(1 * 0.5) + (1 * 1.0) + (0 * 1.2) = 1.5$$



# Activation Function - Threshold Function



$$y = \begin{cases} 1 & \text{if } x \geq \text{threshold} \\ 0 & \text{if } x < \text{threshold} \end{cases}$$

If the value is greater than threshold  $\theta$ , the out put is 1



---

# Class Exercise 10-12 minutes

---


- ❖ Create a threshold logic unit for:
  - ❖ AND Gate
  - ❖ OR Gate



Answer Next Slide



# McCulloch-Pitts two input gates




Threshold for AND: 2  
Threshold for OR: 1

NOTE: No weights



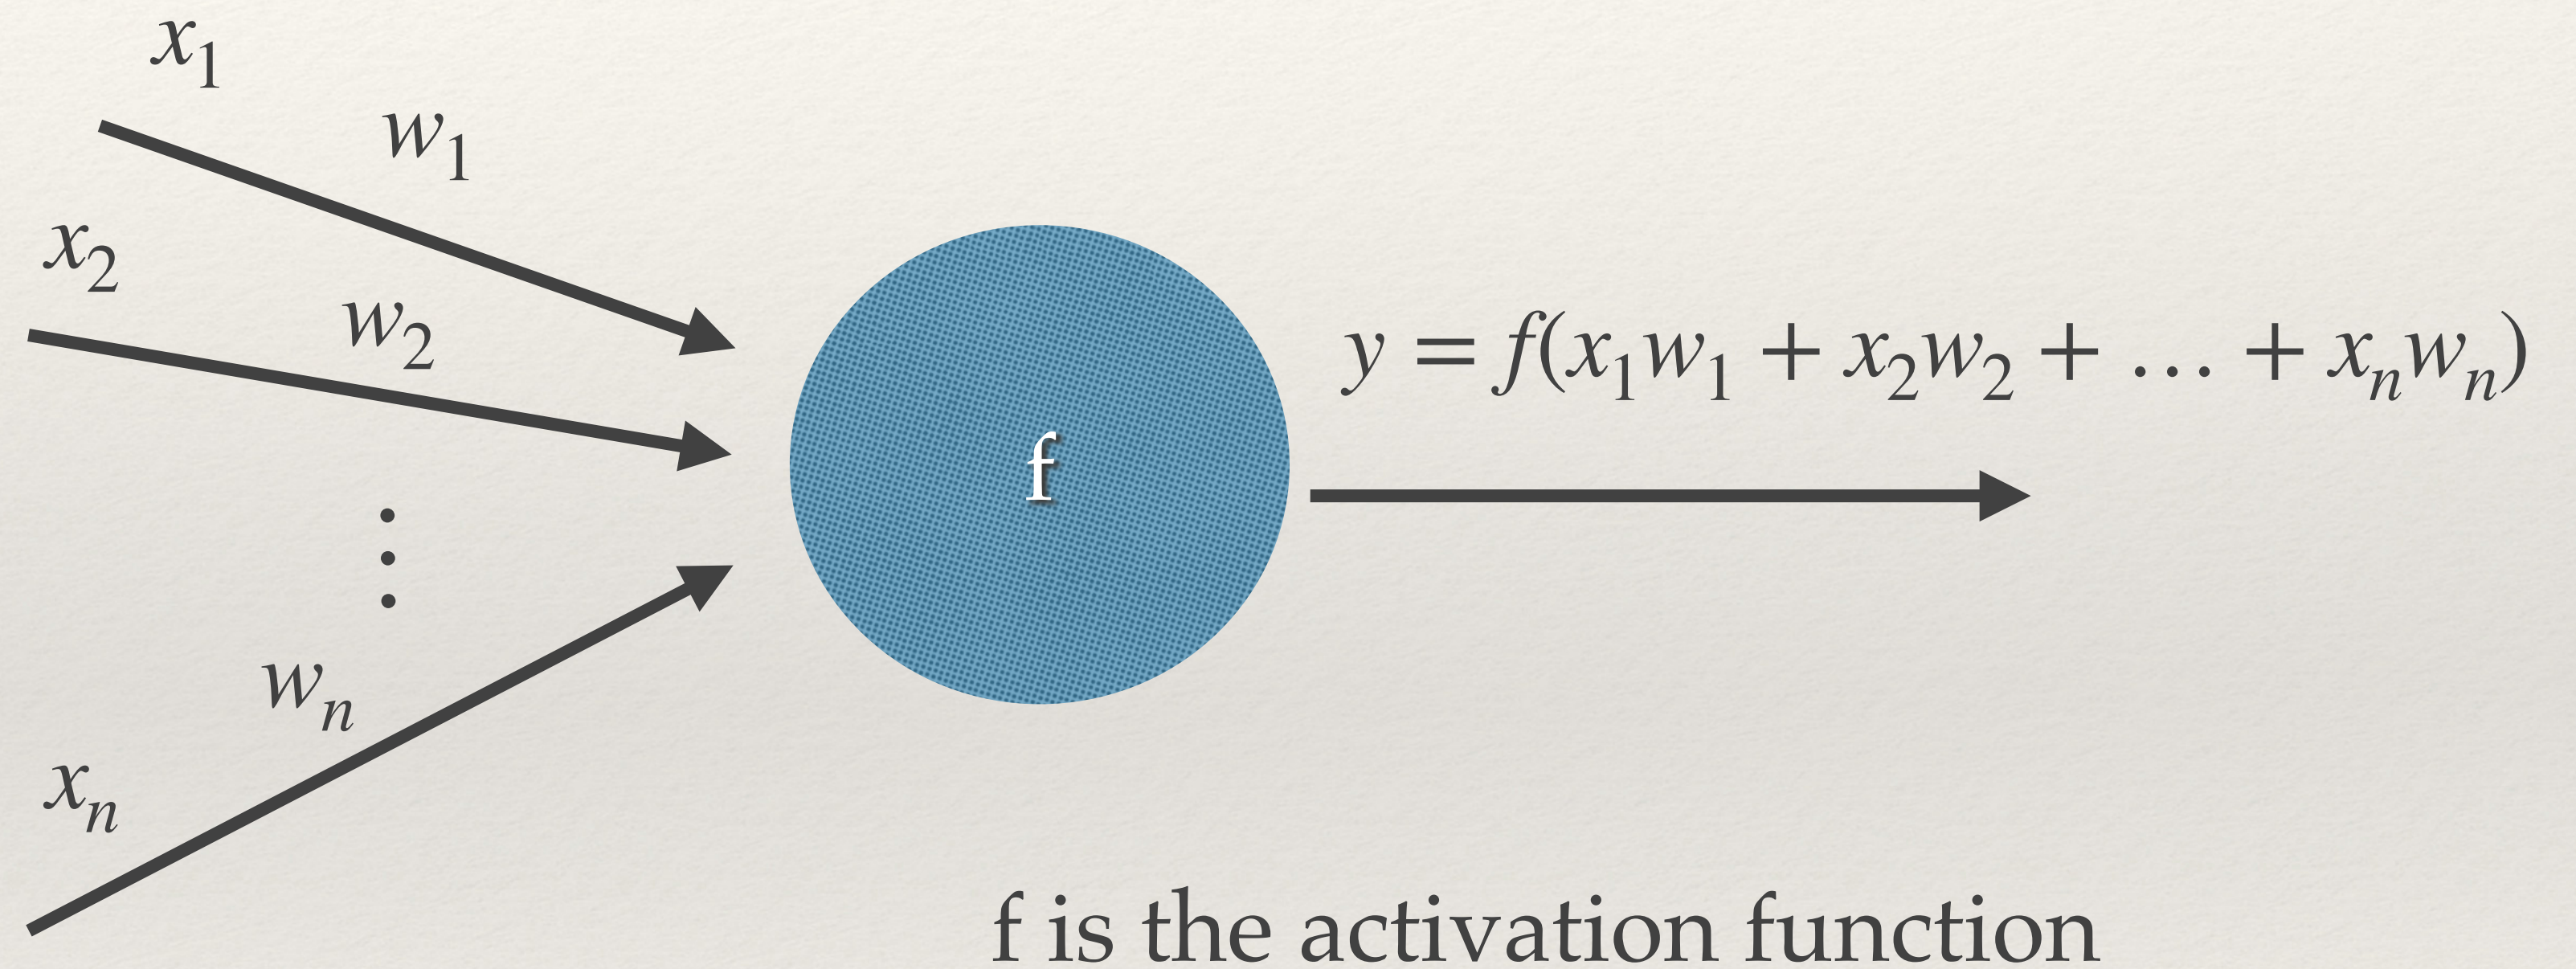
# Neuron






# Activation Functions

- ❖ Sigmoid
- ❖ Tanh
- ❖ ReLu





# Activation Functions - Sigmoid Function



<https://towardsdatascience.com/derivative-of-the-sigmoid-function-536880cf918e>



# Sigmoid Activation Function Properties

- ❖ Range: Output of the sigmoid function is contained between 0 and 1
- ❖ Monotonicity: as the input increases, the output increases also
- ❖ Differentiability: Smooth and differentiable at all points. This is useful for gradient based optimization methods like backpropagation
- ❖ Derivative:  $\sigma'(x) = \sigma(x) \cdot (1 - \sigma(x))$

$$\text{Sigmoid}(x) = \frac{1}{1 + e^{-x}}$$




# Sigmoid Activation Function Properties

- ❖ Gradient Saturation: For large positive or large negative inputs, the sigmoid function saturates. This can lead to the vanishing gradient problem where update to weights become very small slowing down or stopping the learning process
- ❖ Output bound: Sigmoid outputs values close to 0 or 1, which makes it useful in binary classification where we want probabilities as outputs
- ❖ Non-zero-centered: The outputs are always positive, which can sometimes slow down training

$$\text{Sigmoid}(x) = \frac{1}{1 + e^{-x}}$$



# Activation Functions - Tanh Function



<https://towardsdatascience.com/derivative-of-the-sigmoid-function-536880cf918e>



---

# Tanh Activation Function Properties

---

- ❖ Range: Output of the tanh function is contained between -1 and 1, making it zero-centered
- ❖ Zero-centered: Gradient updates are more balanced, leading to more efficient training as weight updates can flow both positively and negatively.
- ❖ S-shaped curve: Smoothly transitions between the bounds -1 and 1
- ❖ Monotonicity: as the input increases, the output increases also
- ❖ Differentiability: Smooth and differentiable at all points. This is useful for gradient based optimization methods like backpropagation
- ❖ Derivative:  $\tanh'(x) = 1 - \tanh^2(x)$




# Tanh Activation Function Properties

- ❖ Symmetry: The tanh function is odd-symmetric about the origin. This means for any input  $x$ ,  $\tanh(-x) = -\tanh(x)$ . This symmetry around zero helps avoid bias in the activation
- ❖ Gradient Saturation: For large positive or large negative inputs, the tanh function saturates. This can lead to the vanishing gradient problem where update to weights become very small slowing down or stopping the learning process
- ❖ Output bound: Sigmoid outputs values close to -1 or 1, tanh is effective in normalizing data and limiting the range of activation values, which can help stabilize learning in certain cases
- ❖ Better for Hidden Layers: Compared to the sigmoid function because of its zero centered output, however in modern deep learning architectures ReLU and its variants are often used instead due to their ability to mitigate the vanishing gradient problem



# Activation Functions - Relu Function



<https://towardsdatascience.com/derivative-of-the-sigmoid-function-536880cf918e>



---

# ReLU Activation Function Properties

---

- ❖ Range: The ReLU function outputs the values in the range of
- ❖ Sparsity: One of ReLU's defining features is that it outputs 0 for any negative input. This leads to sparsity in the activations, meaning any neurons can have zero output at any given time. This sparsity can improve computational efficiency and reduce the complexity of the model, especially in large neural networks.
- ❖ Non-linearity: introduces non-linearity into the model, which is essential for learning complex patterns.
- ❖ Efficient gradient propagation: ReLU alleviates the vanishing gradient problem faced by sigmoid and tang functions. For positive input values, the gradient of ReLU is always 1, ensuring that gradients can flow through the network efficiently. This helps in faster convergence during training, especially in deep networks



---

# ReLU Activation Function Properties

---

- ❖ Derivative: The derivative of ReLU is simple
- ❖ Computational efficiency: it only involves a thresholding operation, making it very fast to compute unlike sigmoid and tang which involve exponentiation
- ❖ Not zero-centered: like the sigmoid function, ReLU is not zero-centered, meaning its output is always non-negative. This can potentially lead to some inefficiencies during training, but this issue is generally outweighed by its other advantages
- ❖ Dying ReLU problem: Key drawback of ReLU, if a neuron receives only negative inputs during training, it will output 0, and its gradient will also be 0, preventing a neuron from learning. Once the neuron can get stuck in this state and never activate again. This can lead to dead neurons that never fire for any input reducing the model's capacity.



---


# ReLU Activation Function Properties (cont'd)

---

- ❖ Variants of ReLU such as Leaky ReLU and Parametric ReLU (PReLU) have been introduced. Leaky ReLU allows a small, non-zero gradient for negative inputs, reducing the chances of neurons dying
- ❖ Used in Hidden layers - more on this to come!
- ❖ Unbounded output: The output of the ReLU can grow indefinitely for positive inputs. While this can be useful for allowing the model to learn larger activations, it can also lead to potential issues like exploding gradients if not properly managed (e.g., using techniques like gradient clipping)



# Activation Functions



<https://towardsdatascience.com/derivative-of-the-sigmoid-function-536880cf918e>



---

# Activation Functions

---

- ❖ Which one should I use?
- ❖ Designing neural networks can feel like more of an 'art' than a 'science'
- ❖ The more we're learning about them, the easier it is to make these decisions
- ❖ <https://deeppai.org/machine-learning-glossary-and-terms/sigmoid-function>



# TLU with two inputs

By setting  $\bar{x} \cdot \bar{w} = \theta$  we obtain:  $x_1 * w_1 + x_2 * w_2 = \theta$ . We solve for  $x_2$  in terms of  $x_1, w_1, w_2$  and  $\theta$ , yielding:

$$x_2 * w_2 = \theta - x_1 * w_1.$$

Some algebraic manipulation gives us:

$$x_2 * w_2 = -x_1 * w_1 + \theta$$

$$x_2 = -\frac{w_1}{w_2} * x_1 + \frac{\theta}{w_2}.$$




# TLU with two inputs

Recall that the equation of a straight line is:  $y = m * x + b$  where  $m$  is the slope ( $m$  equals  $\frac{\Delta y}{\Delta x}$  or the change in  $y$  divided by the change in  $x$ ), and  $b$  is the  $y$ -intercept. Hence, we have a straight line whose slope equals  $-\frac{w_1}{w_2}$  and intercept equals  $\frac{\theta}{w_2}$ . Substituting the values for  $w_1$ ,  $w_2$ ,  $\theta$ , shown in Figure 11.13 we have:  $x_2 = -x_1 + 2$ . This line is shown in Figure 11.14.



# TLU with two inputs



**Figure 11.14**

The straight line obtained from the TLU in Figure 11.13 when the excitation equals the threshold.



# Input/Output Behavior for the TLU

Where  $x_1$  and  $x_2$   
are the input and  $y$   
is the output

$x_1$	$x_2$	$\bar{x} \cdot \bar{w}$	$y$
0	0		
0	1		
1	0		
1	1		

Fill in the columns  
Each weight = 0.5

**Figure 11.15**  
Input/output behavior for the TLU in Figure 11.13.



# Input/Output Behavior for the TLU

Where  $x_1$  and  $x_2$   
are the input and  $y$   
is the output

$x_1$	$x_2$	$\bar{x} \cdot \bar{w}$	$y$
0	0	0.0	0
0	1	0.5	0
1	0	0.5	0
1	1	1.0	1

**Figure 11.15**  
Input/output behavior for the TLU in Figure 11.13.



## Perceptron Learning Algorithm Pseudocode

Inputs:

X: set of input patterns  $\{x_1, x_2, \dots, x_p\}$

D: set of desired outputs  $\{d_1, d_2, \dots, d_p\}$

W: augmented weight vector (randomly initialized)

Algorithm:

// Initialize a variable to check if weights need further updates

not\_all\_equal = true

// Loop until the algorithm converges

while (not\_all\_equal) do:

    not\_all\_equal = false

    // Loop over all input patterns

    for  $i = 1$  to  $p$  do:

        // Compute the perceptron output (using the current weights)

$y_i = \text{thetaActivation}(W \cdot x_i)$  // returns 1 if  $\geq \text{theta}$ , 0 if  $< \text{theta}$

        // If the output is not equal to the desired output

        if ( $y_i \neq d_i$ ) then:

            // Update the weight vector:  $W = W + \Delta W$

$W = W + \eta * (d_i - y_i) * x_i$  //  $\eta$  is the learning rate

            // Set flag to true to continue training

            not\_all\_equal = true

// If no weights were updated in the loop, the model is trained

return "TLU has successfully been trained"



END